

# Class17: Vaccination Rate Mini Project

Wenxi Tang

## ##Background

In today's class, we will explore a dataset on state wide vaccination rate from CA.gov.

The goal of this hands-on mini-project is to examine and compare the Covid-19 vaccination rates around San Diego.

We will start by downloading the most recently dated "Statewide COVID-19 Vaccines Administered by ZIP Code" CSV file from: <https://data.ca.gov/dataset/covid-19-vaccine-progress-dashboard-data-by-zip-code>

## Data input

```
# Import vaccination data
vax <- read.csv("covid19vaccinesbyzipcode_test.csv")
head(vax)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2021-01-05	93609	Fresno	Fresno
2	2021-01-05	94086	Santa Clara	Santa Clara
3	2021-01-05	94304	Santa Clara	Santa Clara
4	2021-01-05	94110	San Francisco	San Francisco
5	2021-01-05	93420	San Luis Obispo	San Luis Obispo
6	2021-01-05	93454	Santa Barbara	Santa Barbara
	vaccine_equity_metric_quartile		vem_source	
1		1	Healthy Places Index Score	
2		4	Healthy Places Index Score	
3		4	Healthy Places Index Score	
4		4	Healthy Places Index Score	
5		3	Healthy Places Index Score	
6		2	Healthy Places Index Score	

	age12_plus_population	age5_plus_population	tot_population
1	4396.3	4839	5177
2	42696.0	46412	50477
3	3263.5	3576	3852
4	64350.7	68320	72380
5	26694.9	29253	30740
6	32043.4	36446	40432

	persons_fully_vaccinated	persons_partially_vaccinated
1	NA	NA
2	11	640
3	NA	NA
4	18	1262
5	NA	NA
6	NA	NA

	percent_of_population_fully_vaccinated
1	NA
2	0.000218
3	NA
4	0.000249
5	NA
6	NA

	percent_of_population_partially_vaccinated
1	NA
2	0.012679
3	NA
4	0.017436
5	NA
6	NA

	percent_of_population_with_1_plus_dose	booster_recip_count
1	NA	NA
2	0.012897	NA
3	NA	NA
4	0.017685	NA
5	NA	NA
6	NA	NA

	bivalent_dose_recip_count	eligible_recipient_count
1	NA	1
2	NA	11
3	NA	6
4	NA	18
5	NA	4
6	NA	5

redacted

1 Information redacted in accordance with CA state privacy requirements  
 2 Information redacted in accordance with CA state privacy requirements  
 3 Information redacted in accordance with CA state privacy requirements  
 4 Information redacted in accordance with CA state privacy requirements  
 5 Information redacted in accordance with CA state privacy requirements  
 6 Information redacted in accordance with CA state privacy requirements

```
library(skimr)
```

A useful function for exploring new datasets is **skimr** package.

```
skimr::skim(vax)
```

Table 1: Data summary

Name	vax
Number of rows	201096
Number of columns	18
Column type frequency:	
character	5
numeric	13
Group variables	None

### Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
as_of_date	0	1	10	10	0	114	0
local_health_jurisdiction	0	1	0	15	570	62	0
county	0	1	0	15	570	59	0
vem_source	0	1	15	26	0	3	0
redacted	0	1	2	69	0	2	0

### Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
zip_code_tabulation_area	0	1.00	93665.11	1817.38	000019	2257.75	3658.50	5380.50	7635.0	
vaccine_equity_metric_quartile	0	0.95	2.44	1.11	1	1.00	2.00	3.00	4.0	

skim_variable	n_missing	complete	mean	sd	p0	p25	p50	p75	p100	hist
age12_plus_population	0	1.00	18895.04	8993.87	0	1346.95	13685.13	1756.18	8556.7	
age5_plus_population	0	1.00	20875.24	1105.97	0	1460.50	15364.08	14877.00	1902.0	
tot_population	9804	0.95	23372.72	2628.50	12	2126.00	18714.08	168.00	11165.0	
persons_fully_vaccinated	16621	0.92	13990.39	5073.66	11	932.00	8589.00	23346.08	7575.0	
persons_partially_vaccinated	16621	0.92	1702.31	2033.32	11	165.00	1197.00	2536.00	39973.0	
percent_of_population_20165_vaccinated	20965	0.90	0.57	0.25	0	0.42	0.61	0.74	1.0	
percent_of_population_20165_fully_vaccinated	20965	0.90	0.08	0.09	0	0.05	0.06	0.08	1.0	
percent_of_population_20165_1_plus_dose	20965	0.89	0.63	0.24	0	0.49	0.67	0.81	1.0	
booster_recip_count	72997	0.64	5882.76	219.00	11	300.00	2773.00	510.00	59593.0	
bivalent_dose_recip_count	158776	0.21	2978.23	3633.03	11	193.00	1467.50	1730.25	27694.0	
eligible_recipient_count	0	1.00	12830.83	4928.64	0	507.00	6369.00	22014.08	7248.0	

```
#sort(vax$as_of_date, decreasing = F)
#sort(vax$as_of_date)
```

Q1. What column details the total number of people fully vaccinated?

The column “persons\_fully\_vaccinated”

Q2. What column details the Zip code tabulation area?

The column “zip\_code\_tabulation\_area”

Q3. What is the earliest date in this dataset?

```
vax$as_of_date[1]
```

```
[1] "2021-01-05"
```

2021-01-05

Q4. What is the latest date in this dataset?

```
vax$as_of_date[nrow(vax)]
```

```
[1] "2023-03-07"
```

2023-03-07

Q5. How many numeric columns are in this dataset?

13

Q6. Note that there are “missing values” in the dataset. How many NA values there in the `persons_fully_vaccinated` column?

```
sum(is.na(vax$persons_fully_vaccinated))
```

```
[1] 16621
```

16621

Q7. What percent of `persons_fully_vaccinated` values are missing (to 2 significant figures)?

```
round(16621/nrow(vax) * 100, 2)
```

```
[1] 8.27
```

Q8. [Optional]: Why might this data be missing?

It might be some confidential federal areas, or at very early stages, data was not collected.

##Working with dates

We will use the *lubridate* package to help ease the pain of working with times and dates.

```
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

```
date, intersect, setdiff, union
```

```
today()
```

```
[1] "2023-03-08"
```

Q9. How many days have passed since the last update of the dataset?

```
today() - ymd(vax$as_of_date[nrow(vax)])
```

Time difference of 1 days

I will convert the entire “as\_of\_date” column to be in lubridate format.

```
vax$as_of_date <- ymd(vax$as_of_date)
```

```
today() - vax$as_of_date[(nrow(vax))]
```

Time difference of 1 days

```
today() - ymd("1997-4-20")
```

Time difference of 9453 days

Q10. How many unique dates are in the dataset (i.e. how many different dates are detailed)?

```
length(unique(vax$as_of_date))
```

```
[1] 114
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
n_distinct(vax$as_of_date)
```

```
[1] 114
```

## Working with zipcode

There are quite a few R packages that can help ease the pain of working with ZIP codes, we will try just one of the smaller ones **zipcodeR**.

```
library(zipcodeR)
```

```
geocode_zip('92037')
```

```
# A tibble: 1 x 3
  zipcode lat lng
  <chr>   <dbl> <dbl>
1 92037   32.8 -117.
```

```
zip_distance('92037','92109')
```

```
zipcode_a zipcode_b distance
1      92037      92109      2.33
```

```
#Calculate the distance between the centroids of any two ZIP codes in miles, e.g.
reverse_zipcode(c('92037', "92109", "92122", "94061")) )
```

```
# A tibble: 4 x 24
  zipcode zipcode_~1 major~2 post_~3 common_c~4 county state lat lng timez~5
  <chr>   <chr>      <chr>   <chr>      <blob> <chr>   <chr> <dbl> <dbl> <chr>
1 92037   Standard    La Jol~ La Jol~ <raw 20 B> San D~ CA    32.8 -117. Pacific
2 92109   Standard    San Di~ San Di~ <raw 21 B> San D~ CA    32.8 -117. Pacific
3 92122   Standard    San Di~ San Di~ <raw 21 B> San D~ CA    32.9 -117. Pacific
4 94061   Standard    Redwoo~ Redwoo~ <raw 34 B> San M~ CA    37.5 -122. Pacific
# ... with 14 more variables: radius_in_miles <dbl>, area_code_list <blob>,
#   population <int>, population_density <dbl>, land_area_in_sqmi <dbl>,
#   water_area_in_sqmi <dbl>, housing_units <int>,
#   occupied_housing_units <int>, median_home_value <int>,
```

```
# median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
# bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
# 1: zipcode_type, 2: major_city, 3: post_office_city, ...
```

Q. Find the best and worst ratio of “median household income” to “median\_home\_value” in San Diego county

Q11. How many distinct zip codes are listed for San Diego County?

```
sd_zip <- unique(vax$zip_code_tabulation_area[vax$county == "San Diego"])
length(sd_zip)
```

```
[1] 107
```

```
sd.eco <- reverse_zipcode(sd_zip)
```

Now extract the best and worst “median household income” and “median\_home\_value”

```
sort(sd.eco$median_household_income/sd.eco$median_home_value, decreasing = T)
```

```
[1] 0.49313299 0.35517505 0.34548798 0.28576236 0.25878896 0.25671904
[7] 0.25053447 0.22687477 0.22497633 0.22413397 0.22139649 0.21978362
[13] 0.21812314 0.21737751 0.21539376 0.21355819 0.21243233 0.21139772
[19] 0.20679398 0.20651554 0.20551950 0.20209350 0.20191970 0.19888657
[25] 0.19792504 0.19422057 0.18908834 0.18908362 0.18833910 0.18786741
[31] 0.18769501 0.18727151 0.18673482 0.18672321 0.18607322 0.18288752
[37] 0.18134523 0.18028781 0.17915189 0.17828642 0.17770952 0.17550646
[43] 0.17179217 0.17147450 0.16887196 0.16562207 0.16401408 0.16314344
[49] 0.15958461 0.15884808 0.15613789 0.15532406 0.15500482 0.15364149
[55] 0.15001161 0.14930357 0.14815293 0.14725724 0.14711029 0.14562657
[61] 0.14452578 0.14452332 0.14416319 0.14416195 0.14400708 0.14313883
[67] 0.14243254 0.14100882 0.14073491 0.14019337 0.13936012 0.13388495
[73] 0.13376211 0.13327181 0.13265276 0.12806991 0.12505087 0.12381000
[79] 0.12324040 0.12138591 0.11951699 0.11711611 0.11622243 0.11618388
[85] 0.11487989 0.11445052 0.11431968 0.11381931 0.11355523 0.11223549
[91] 0.10313557 0.09936690 0.09831696 0.09363608 0.09228770 0.09110291
[97] 0.08988986
```

```
arrange(sd.eco, sd.eco$median_household_income, sd.eco$median_home_value)
```



```
# A tibble: 107 x 24
  zipcode zipcode~1 major~2 post_~3 common_c~4 county state lat lng timez~5
  <chr> <chr> <chr> <chr> <blob> <chr> <chr> <dbl> <dbl> <chr>
1 91931 PO Box Guatay Guatay~ <raw 18 B> San D~ CA 32.8 -117. Pacific
2 92113 Standard San Di~ San Di~ <raw 21 B> San D~ CA 32.7 -117. Pacific
3 92173 Standard San Ys~ San Ys~ <raw 31 B> San D~ CA 32.6 -117. Pacific
4 92066 Standard Ranchi~ Ranchi~ <raw 40 B> San D~ CA 33.3 -116. Pacific
5 92105 Standard San Di~ San Di~ <raw 21 B> San D~ CA 32.7 -117. Pacific
6 91950 Standard Nation~ Nation~ <raw 25 B> San D~ CA 32.7 -117. Pacific
7 92086 Standard Warner~ Warner~ <raw 31 B> San D~ CA 33.3 -117. Pacific
8 91963 Standard Potrero Potrer~ <raw 19 B> San D~ CA 32.6 -117. Pacific
9 92004 Standard Borreg~ Borreg~ <raw 32 B> San D~ CA 33.1 -116. Pacific
10 92102 Standard San Di~ San Di~ <raw 21 B> San D~ CA 32.7 -117. Pacific
# ... with 97 more rows, 14 more variables: radius_in_miles <dbl>,
# area_code_list <blob>, population <int>, population_density <dbl>,
# land_area_in_sqmi <dbl>, water_area_in_sqmi <dbl>, housing_units <int>,
# occupied_housing_units <int>, median_home_value <int>,
# median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
# bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
# 1: zipcode_type, 2: major_city, 3: post_office_city, ...
```

```
# Pull data for all ZIP codes in the dataset
#zipdata <- reverse_zipcode( vax$zip_code_tabulation_area )
```

## Focus on the San Diego area

Q12. What San Diego County Zip code area has the largest 12 + Population in this dataset?

```
sd <- filter(vax, county == "San Diego")
```

```
ind <- which.max(sd$age12_plus_population)
sd[ind,]
```

```
as_of_date zip_code_tabulation_area local_health_jurisdiction county
67 2021-01-05 92154 San Diego San Diego
vaccine_equity_metric_quartile vem_source
67 2 Healthy Places Index Score
age12_plus_population age5_plus_population tot_population
67 76365.2 82971 88979
```

```

persons_fully_vaccinated persons_partially_vaccinated
67                      16                      1400
percent_of_population_fully_vaccinated
67                      0.00018
percent_of_population_partially_vaccinated
67                      0.015734
percent_of_population_with_1_plus_dose booster_recip_count
67                      0.015914                      NA
bivalent_dose_recip_count eligible_recipient_count
67                      NA                      16
redacted
67 Information redacted in accordance with CA state privacy requirements

```

Q13. What is the overall average “Percent of Population Fully Vaccinated” value for all San Diego “County” as of “2023-03-07”?

```
vax$as_of_date[nrow(vax)]
```

```
[1] "2023-03-07"
```

```
thisdate <- filter(sd, as_of_date == "2023-03-07")
```

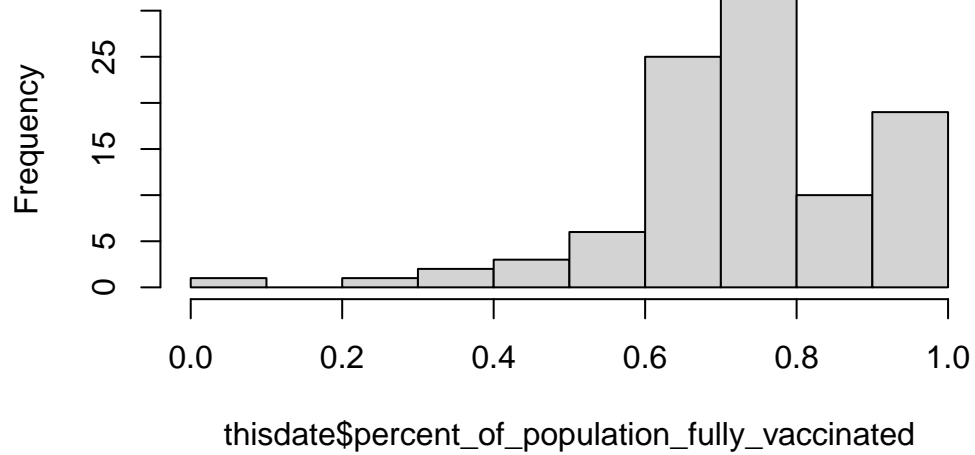
```
mean(thisdate$percent_of_population_fully_vaccinated, na.rm = T)
```

```
[1] 0.7402567
```

Q14. Using either ggplot or base R graphics make a summary figure that shows the distribution of Percent of Population Fully Vaccinated values as of “2023-02-28”?

```
hist(thisdate$percent_of_population_fully_vaccinated)
```


## Histogram of thisdate\$percent\_of\_population\_fully\_vaccinated



```
library(ggplot2)

ggplot(thisdate) +
  aes(percent_of_population_fully_vaccinated) +
  geom_histogram(binwidth = 0.1) +
  theme_bw() +
  xlab("percent of population fully vaccinated") +
  ylab("Counts")
```

Warning: Removed 8 rows containing non-finite values (`stat\_bin()`).



### Focus on UCSD/La Jolla

UCSD resides in the 92037 zip code area.


```
ucsd <- filter(sd, zip_code_tabulation_area=="92037")
ucsd[1,]$age5_plus_population
```

[1] 36144

Q15. Using ggplot make a graph of the vaccination rate time course for the 92037 ZIP code area:

```
ljplot <- ggplot(ucsd) +
  aes(as_of_date,
      percent_of_population_fully_vaccinated) +
  geom_point() +
  geom_line(group=1) +
  ylim(c(0,1)) +
  labs(x = "Date", y="Percent Vaccinated")
```

ljplot



## Comparing to similar sized areas

Let's return to the full dataset and look across every zip code area with a population at least as large as that of 92037 on as\_of\_date "2023-03-07".

```
# Subset to all CA areas with a population as large as 92037
vax.36 <- filter(vax, age5_plus_population > 36144 &
  as_of_date == "2023-03-07")

head(vax.36)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2023-03-07	94116	San Francisco	San Francisco
2	2023-03-07	92703	Orange	Orange
3	2023-03-07	94118	San Francisco	San Francisco
4	2023-03-07	92376	San Bernardino	San Bernardino
5	2023-03-07	92692	Orange	Orange
6	2023-03-07	95148	Santa Clara	Santa Clara

	vaccine_equity_metric_quartile	vem_source
1	4	Healthy Places Index Score
2	1	Healthy Places Index Score
3	4	Healthy Places Index Score

4	1 Healthy Places Index Score		
5	4 Healthy Places Index Score		
6	4 Healthy Places Index Score		
	age12_plus_population	age5_plus_population	tot_population
1	42334.3	45160	47346
2	57182.7	64387	69112
3	37628.5	40012	42095
4	70232.1	79686	86085
5	41008.9	44243	46800
6	42163.3	46202	48273
	persons_fully_vaccinated	persons_partially_vaccinated	
1	41255	2450	
2	57887	7399	
3	33284	3040	
4	51367	5674	
5	35117	2603	
6	42298	2684	
	percent_of_population_fully_vaccinated		
1	0.871351		
2	0.837582		
3	0.790688		
4	0.596701		
5	0.750363		
6	0.876225		
	percent_of_population_partially_vaccinated		
1	0.051747		
2	0.107058		
3	0.072218		
4	0.065912		
5	0.055620		
6	0.055600		
	percent_of_population_with_1_plus_dose	booster_recip_count	
1	0.923098	34108	
2	0.944640	28297	
3	0.862906	27401	
4	0.662613	23832	
5	0.805983	23695	
6	0.931825	31583	
	bivalent_dose_recip_count	eligible_recipient_count	redacted
1	19158	41000	No
2	7627	57775	No
3	15251	33146	No
4	6393	51276	No

5	10169	35031	No
6	12604	42120	No

Q16. Calculate the mean “Percent of Population Fully Vaccinated” for ZIP code areas with a population as large as 92037 (La Jolla) as\_of\_date “2023-02-28”. Add this as a straight horizontal line to your plot from above with the `geom_hline()` function?


```
n_distinct(vax.36$zip_code_tabulation_area)
```

```
[1] 411
```

```
mean(vax.36$percent_of_population_fully_vaccinated, na.rm = T)
```

```
[1] 0.7214936
```

```
ljplot +  
  geom_hline(yintercept = 0.7214936, col = "red", linetype = 5)
```



Q17. What is the 6 number summary (Min, 1st Qu., Median, Mean, 3rd Qu., and Max) of the “Percent of Population Fully Vaccinated” values for ZIP code areas with a population as large as 92037 (La Jolla) as\_of\_date “2023-03-07”?


```
summary(vax.36$percent_of_population_fully_vaccinated)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.3805	0.6459	0.7183	0.7215	0.7908	1.0000

Q18. Using ggplot generate a histogram of this data.

```
ggplot(vax.36) +  
  aes(percent_of_population_fully_vaccinated) +  
  geom_histogram(bins = 20) +  
  xlim(0,1) +  
  geom_vline(xintercept = 0.7215, col = "red", linetype = 5)
```

Warning: Removed 2 rows containing missing values (`geom\_bar()`).



Q19. Is the 92109 and 92040 ZIP code areas above or below the average value you calculated for all these above?



```
vax %>% filter(as_of_date == "2023-03-07") %>%
  filter(zip_code_tabulation_area== "92040") %>%
  select(percent_of_population_fully_vaccinated)
```

```
percent_of_population_fully_vaccinated
1                                0.550533
```

```
filter(vax.36, zip_code_tabulation_area %in% c("92109", "92040"))
```

```
as_of_date zip_code_tabulation_area local_health_jurisdiction county
1 2023-03-07                92109                San Diego San Diego
2 2023-03-07                92040                San Diego San Diego
vaccine_equity_metric_quartile          vem_source
1                                3 Healthy Places Index Score
2                                3 Healthy Places Index Score
age12_plus_population age5_plus_population tot_population
1                43222.5                44953                47111
2                39405.0                42833                46306
persons_fully_vaccinated persons_partially_vaccinated
1                32725                4234
2                25493                2156
percent_of_population_fully_vaccinated
1                0.694636
2                0.550533
percent_of_population_partially_vaccinated
1                0.089873
2                0.046560
percent_of_population_with_1_plus_dose booster_recip_count
1                0.784509                19677
2                0.597093                14175
bivalent_dose_recip_count eligible_recipient_count redacted
1                8109                32622        No
2                4649                25433        No
```

They are both below 0.7215.

Q20. Finally make a time course plot of vaccination progress for all areas in the full dataset with a age5\_plus\_population > 36144.

```
last <- vax[vax$age5_plus_population > 36144,]
head(last)
```

	as_of_date	zip_code	tabulation_area	local_health_jurisdiction	county									
2	2021-01-05		94086	Santa Clara	Santa Clara									
4	2021-01-05		94110	San Francisco	San Francisco									
6	2021-01-05		93454	Santa Barbara	Santa Barbara									
7	2021-01-05		93458	Santa Barbara	Santa Barbara									
9	2021-01-05		93536	Los Angeles	Los Angeles									
12	2021-01-05		94066	San Mateo	San Mateo									
	vaccine_equity_metric_quartile			vem_source										
2	4 Healthy Places			Index Score										
4	4 Healthy Places			Index Score										
6	2 Healthy Places			Index Score										
7	1 Healthy Places			Index Score										
9	2 Healthy Places			Index Score										
12	4 Healthy Places			Index Score										
	age12_plus_population	age5_plus_population	tot_population											
2	42696.0	46412	50477											
4	64350.7	68320	72380											
6	32043.4	36446	40432											
7	43262.5	51006	57256											
9	59659.1	65612	70237											
12	37730.3	40903	43101											
	persons_fully_vaccinated	persons_partially_vaccinated												
2	11	640												
4	18	1262												
6	NA	NA												
7	NA	NA												
9	53	1066												
12	NA	NA												
	percent_of_population_fully_vaccinated													
2	0.000218													
4	0.000249													
6	NA													
7	NA													
9	0.000755													
12	NA													
	percent_of_population_partially_vaccinated													
2	0.012679													
4	0.017436													
6	NA													

7		NA
9		0.015177
12		NA
	percent_of_population_with_1_plus_dose	booster_recip_count
2		0.012897 NA
4		0.017685 NA
6		NA NA
7		NA NA
9		0.015932 NA
12		NA NA
	bivalent_dose_recip_count	eligible_recipient_count
2	NA	11
4	NA	18
6	NA	5
7	NA	2
9	NA	53
12	NA	3
		redacted
2	Information redacted in accordance with CA state privacy requirements	
4	Information redacted in accordance with CA state privacy requirements	
6	Information redacted in accordance with CA state privacy requirements	
7	Information redacted in accordance with CA state privacy requirements	
9	Information redacted in accordance with CA state privacy requirements	
12	Information redacted in accordance with CA state privacy requirements	


```
vax.36.all <- filter(vax, age5_plus_population > 36144)
```

```
ggplot(vax.36.all) +
  aes(as_of_date,
      percent_of_population_fully_vaccinated,
      group=zip_code_tabulation_area) +
  geom_line(alpha=0.2, color="blue") +
  ylim(0,1) +
  labs(x="Date", y="Percent of Population Fully Vaccination",
       title="Vaccination Rate across California",
       subtitle="Only areas with a population above 36k are shown") +
  geom_hline(yintercept = 0.7215, linetype=5)
```

Warning: Removed 183 rows containing missing values (`geom\_line()`).

## Vaccination Rate across California

Only areas with a population above 36k are shown



Q21. How do you feel about traveling for Spring Break and meeting for in-person class afterwards?

I feel pretty safe because the vaccination rate is pretty high.